

```

In [1]:
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score

from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier

In [2]:
#A IS THE DATAFRAME OF WATER QUALITY DATASET
A=pd.read_csv("C:/Users/dell/Desktop/NANDU/PG/SEM2/PROJECT/WATERQUALITY.csv")
A.head(5)

```

Out[2]:

	ph	Hardness	Solids	Chloramines	Sulfate	Conductivity	Organic_carbon	Trihalomethanes	Turbidity	Potability
0	NaN	204.890455	20791.318981	7.300212	368.516441	564.308654	10.379783	86.990970	2.963135	0
1	3.716080	129.422921	18630.057858	6.635246	NaN	592.885359	15.180013	56.329076	4.500656	0
2	8.099124	224.236259	19909.541732	9.275884	NaN	418.606213	16.868637	66.420093	3.055934	0
3	8.316766	214.373394	22018.417441	8.059332	356.886136	363.266516	18.436524	100.341674	4.628771	0
4	9.092223	181.101509	17978.986339	6.546600	310.135738	398.410813	11.558279	31.997993	4.075075	0

```

In [3]:
A.info()
<class 'pandas.core.frame.DataFrame'>

```

RangeIndex: 3276 entries, 0 to 3275

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	ph	2785 non-null	float64
1	Hardness	3276 non-null	float64
2	Solids	3276 non-null	float64
3	Chloramines	3276 non-null	float64
4	Sulfate	2495 non-null	float64
5	Conductivity	3276 non-null	float64
6	Organic_carbon	3276 non-null	float64
7	Trihalomethanes	3114 non-null	float64
8	Turbidity	3276 non-null	float64
9	Potability	3276 non-null	int64

dtypes: float64(9), int64(1)

memory usage: 256.1 KB

In [4]:

#size of A

A.shape

Out[4]:

(3276, 10)

In []:

In [5]:

A.isnull().sum()

Out[5]:

ph	491
Hardness	0
Solids	0
Chloramines	0
Sulfate	781
Conductivity	0
Organic_carbon	0
Trihalomethanes	162
Turbidity	0
Potability	0

dtype: int64

In [6]:

```
A['ph'] = A['ph'].fillna(A.groupby(['Potability'])['ph'].transform('mean'))
```

```
A['Sulfate'] = A['Sulfate'].fillna(A.groupby(['Potability'])['Sulfate'].transform('mean'))
```

```
A['Trihalomethanes'] = A['Trihalomethanes'].fillna(A.groupby(['Potability'])['Trihalomethanes'].transform('mean'))
```

In [7]:

```
A.head(5)
```

Out[7]:

	ph	Hardness	Solids	Chloramines	Sulfate	Conductivity	Organic_carbon	Trihalomethanes	Turbidity	Potability
0	7.085378	204.890455	20791.318981	7.300212	368.516441	564.308654	10.379783	86.990970	2.963135	0
1	3.716080	129.422921	18630.057858	6.635246	334.564290	592.885359	15.180013	56.329076	4.500656	0
2	8.099124	224.236259	19909.541732	9.275884	334.564290	418.606213	16.868637	66.420093	3.055934	0
3	8.316766	214.373394	22018.417441	8.059332	356.886136	363.266516	18.436524	100.341674	4.628771	0
4	9.092223	181.101509	17978.986339	6.546600	310.135738	398.410813	11.558279	31.997993	4.075075	0

In [8]:

```
#CORRELATION BETWEEN VARIABLES
```


```
plt.figure(figsize=(12, 8))
```

```
sns.heatmap(A.corr(), annot=True, cmap="Blues")
```

```
plt.title("Correlations Between Variables", size=16)
```

```
plt.show()
```

Correlations Between Variables



In [9]:

```
#SPLITTING X AND Y
```

```
X = A.drop("Potability", axis=1)
```

```
y = A["Potability"]
```

In [10]:

```
#TEST AND TRAIN DATA SPLIT
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

In [11]:

```
#LOGISTIC REGRESSION MODEL
```

```
LR=LogisticRegression(random_state=42).fit(X_train, y_train)
```

```
y_pred1 = LR.predict(X_test)
```

```
score_LR = accuracy_score(y_test, y_pred1)
```

```
print(score_LR)
```

```
0.6280487804878049
```

In [12]:

```
#DECISION TREE CLASSIFIER MODEL
```

```
DT=DecisionTreeClassifier(random_state=42).fit(X_train, y_train)
```

```
y_pred2=DT.predict(X_test)
```

```
score_DT=accuracy_score(y_test,y_pred2)
```

```
print(score_DT)
```

```
0.7469512195121951
```

In [13]:

```
#KNN model
```

```
KNN=KNeighborsClassifier(n_neighbors=2).fit(X_train,y_train)
```

```
y_pred3=KNN.predict(X_test)
```

```
score_KNN=accuracy_score(y_test,y_pred3)
```

```
print(score_KNN)
```

```
0.5899390243902439
```